

Case study #1: Object counting


Lecturer: Bui Ha Duc, PhD

Email: ducbh@hcmute.edu.vn



Target

- Count the number of target objects in a image



How many cells in this image?



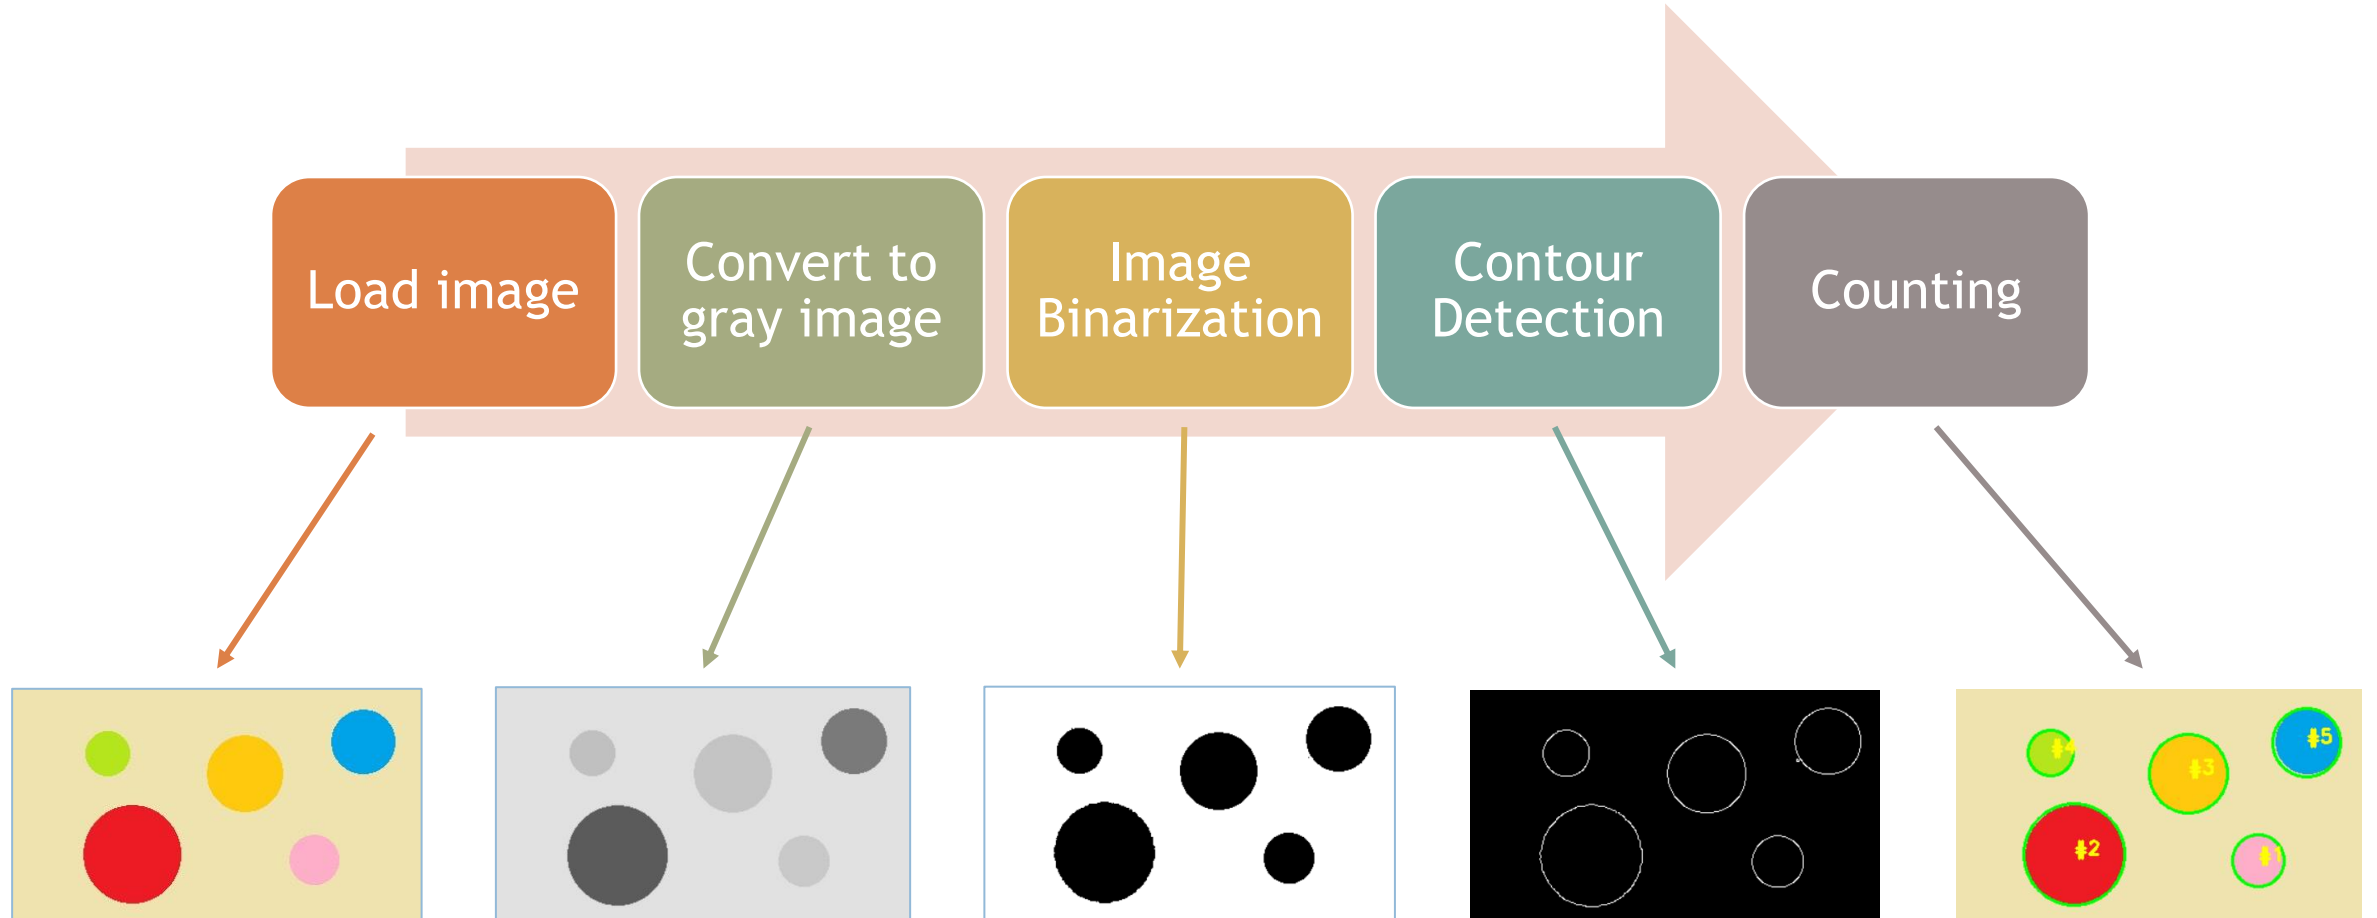
Knowledge will be covered

In this case you will learn

- How to convert an image to gray
- How to binarize a gray image using threshold techniques
- How to find object contours
- How to mark and label object on an image



Simple Object Counting Procedure



Color Conversion


- RGB to Gray

RGB[A] to Gray: $Y \leftarrow 0.299 \cdot R + 0.587 \cdot G + 0.114 \cdot B$


- OpenCV Function:

`cvtColor(src, gray, COLOR_BGR2GRAY);`

Source Image



Gray Image



Conversion mode

- COLOR_BGR2GRAY
- COLOR_BGR2HSV
- COLOR_HSV2BGR
- COLOR_BGR2Lab

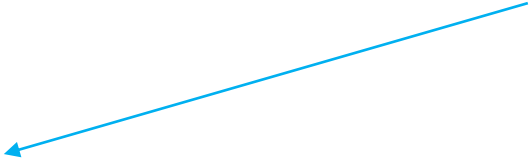
```
// load image file
Mat img = imread(argv[1],1);
// convert image to gray
Mat gray;
cvtColor(img, gray,COLOR_BGR2GRAY);
```



Thresholding

- Purpose: convert a gray image to a binary image
- OpenCV function:

```
threshold(src,dst,threshold_value,max_value,threshold_type);
```



```
threshold_value (THRESH_BINARY):  
if pixel_value < threshold_value  
    pixel_value = 0;  
else  
    pixel_value = max value;
```

```
// convert to binary image using threshold  
Mat b_img;  
double thresh = 115;  
double maxValue = 255;  
  
threshold(gray,b_img,thresh,maxValue,THRESH_BINARY);
```



Thresholding

- Threshold Types

Enumerator	
THRESH_BINARY Python: cv.THRESH_BINARY	$\text{dst}(x, y) = \begin{cases} \text{maxval} & \text{if } \text{src}(x, y) > \text{thresh} \\ 0 & \text{otherwise} \end{cases}$
THRESH_BINARY_INV Python: cv.THRESH_BINARY_INV	$\text{dst}(x, y) = \begin{cases} 0 & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{maxval} & \text{otherwise} \end{cases}$
THRESH_TRUNC Python: cv.THRESH_TRUNC	$\text{dst}(x, y) = \begin{cases} \text{threshold} & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{src}(x, y) & \text{otherwise} \end{cases}$
THRESH_TOZERO Python: cv.THRESH_TOZERO	$\text{dst}(x, y) = \begin{cases} \text{src}(x, y) & \text{if } \text{src}(x, y) > \text{thresh} \\ 0 & \text{otherwise} \end{cases}$
THRESH_TOZERO_INV Python: cv.THRESH_TOZERO_INV	$\text{dst}(x, y) = \begin{cases} 0 & \text{if } \text{src}(x, y) > \text{thresh} \\ \text{src}(x, y) & \text{otherwise} \end{cases}$
THRESH_MASK Python: cv.THRESH_MASK	
THRESH_OTSU Python: cv.THRESH_OTSU	flag, use Otsu algorithm to choose the optimal threshold value
THRESH_TRIANGLE Python: cv.THRESH_TRIANGLE	flag, use Triangle algorithm to choose the optimal threshold value



Contour Detection

▪ Contours

- curve joining all the **continuous points** (along the boundary), having **same color** or **intensity**.
- are a useful tool for **shape analysis**, **object detection** and **recognition**.

▪ OpenCV Function:

findContours (**src**, **contours**, **hierarchy**, **find_mode**, **approx_method**, **offset**)

```
vector<vector<Point> > contours;  
findContours( img, contours, RETR_EXTERNAL, CHAIN_APPROX_SIMPLE );
```

contours

Detected contours.
Each contour is
stored as a vector of
vectors of points.

hierarchy (Optional)

Output vector,
containing
information about the
image topology.

find_mode

Contour
retrieval mode

approx_method

Contour
approximation
method



Contour Detection

- Contour Retrieval Mode

Enumerator	
RETR_EXTERNAL Python: cv.RETR_EXTERNAL	retrieves only the extreme outer contours. It sets <code>hierarchy[i][2]=hierarchy[i][3]=-1</code> for all the contours.
RETR_LIST Python: cv.RETR_LIST	retrieves all of the contours without establishing any hierarchical relationships.
RETR_CCOMP Python: cv.RETR_CCOMP	retrieves all of the contours and organizes them into a two-level hierarchy. At the top level, there are external boundaries of the components. At the second level, there are boundaries of the holes. If there is another contour inside a hole of a connected component, it is still put at the top level.
RETR_TREE Python: cv.RETR_TREE	retrieves all of the contours and reconstructs a full hierarchy of nested contours.
RETR_FLOODFILL Python: cv.RETR_FLOODFILL	




Contour Detection

- Approximation Modes

Enumerator	
CHAIN_APPROX_NONE Python: cv.CHAIN_APPROX_NONE	stores absolutely all the contour points. That is, any 2 subsequent points (x1,y1) and (x2,y2) of the contour will be either horizontal, vertical or diagonal neighbors, that is, $\max(\text{abs}(x1-x2), \text{abs}(y2-y1))=1$.
CHAIN_APPROX_SIMPLE Python: cv.CHAIN_APPROX_SIMPLE	compresses horizontal, vertical, and diagonal segments and leaves only their end points. For example, an up-right rectangular contour is encoded with 4 points.
CHAIN_APPROX_TC89_L1 Python: cv.CHAIN_APPROX_TC89_L1	applies one of the flavors of the Teh-Chin chain approximation algorithm [212]
CHAIN_APPROX_TC89_KCOS Python: cv.CHAIN_APPROX_TC89_KCOS	applies one of the flavors of the Teh-Chin chain approximation algorithm [212]



Problem 1: Image noise



noise

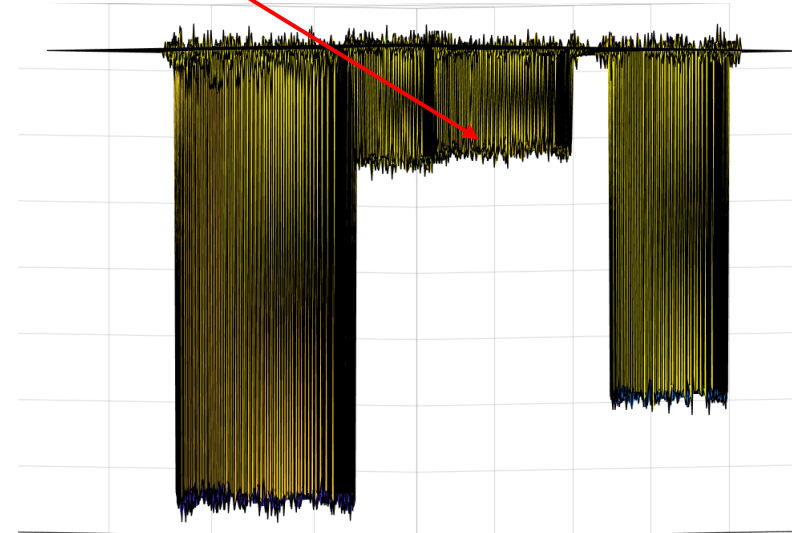

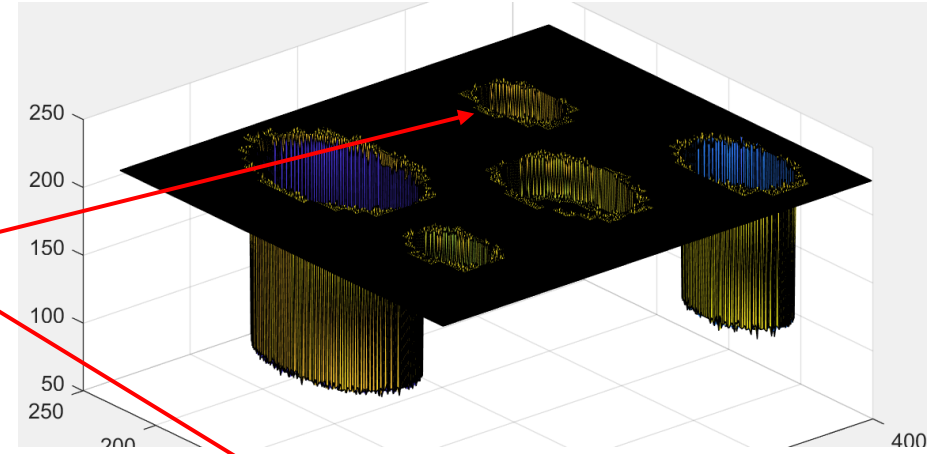


Image smoothing

- Image smoothing: is a key technique in image processing used to reduces noise within an image






Image kernel

- Visual example: <https://setosa.io/ev/image-kernels/>
- Image kernel is a small **matrix** used to apply effects/algorithm.

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

5x5 square

0	0	1	0	0
0	1	1	1	0
1	1	1	1	1
0	1	1	1	0
0	0	1	0	0

5x5 circle

0	0	1	0	0
0	0	1	0	0
1	1	1	1	1
0	0	1	0	0
0	0	1	0	0

5x5 cross



anchor



Kernel multiplication

0	0	0	0	0	0	0
0	105	102	100	97	96	
0	103	99	103	101	102	10
0	101	98	104	102	100	1
0	99	101	106	104	99	
0	104	104	104	100	98	1

Kernel Matrix

0	-1	0
-1	5	-1
0	-1	0

320					

Image Matrix

$$\begin{aligned} &0 * 0 + 0 * -1 + 0 * 0 \\ &+ 0 * -1 + 105 * 5 + 102 * -1 \\ &+ 0 * 0 + 103 * -1 + 99 * 0 = 320 \end{aligned}$$

Output Matrix

Convolution with horizontal and
vertical strides = 2



Image smoothing



- Opencv functions

`blur(src,dst,kernel_size,anchor)` → Mean filter, blur edges
`medianBlur(src,dst,kernel_size)` → Median Filter, preserve edges
`bilateralFilter(src,dst,diameter,sigmaColor,SigmaSpace)`
`GaussianBlur(src,dst,kernel_size,sigmaX,sigmaY)` → Mean filter with Gaussian kernel

```
for(size_t i=1; i<5;i++)
{
    Mat blur_img;
    //medianBlur(gray, blur_img,2*i+1);
    // blur(gray,blur_img,Size(2*i+1,2*i+1), Point(-1,-1));
    GaussianBlur(gray,blur_img,Size(2*i+1,2*i+1),0,0);
    // bilateralFilter(gray,blur_img,i*2,i*2,i*2);
    String window = "kernel = " + to_string(2*i+1) + "x" + to_string(2*i+1);
    imshow( window, blur_img);
}
```



Image smoothing



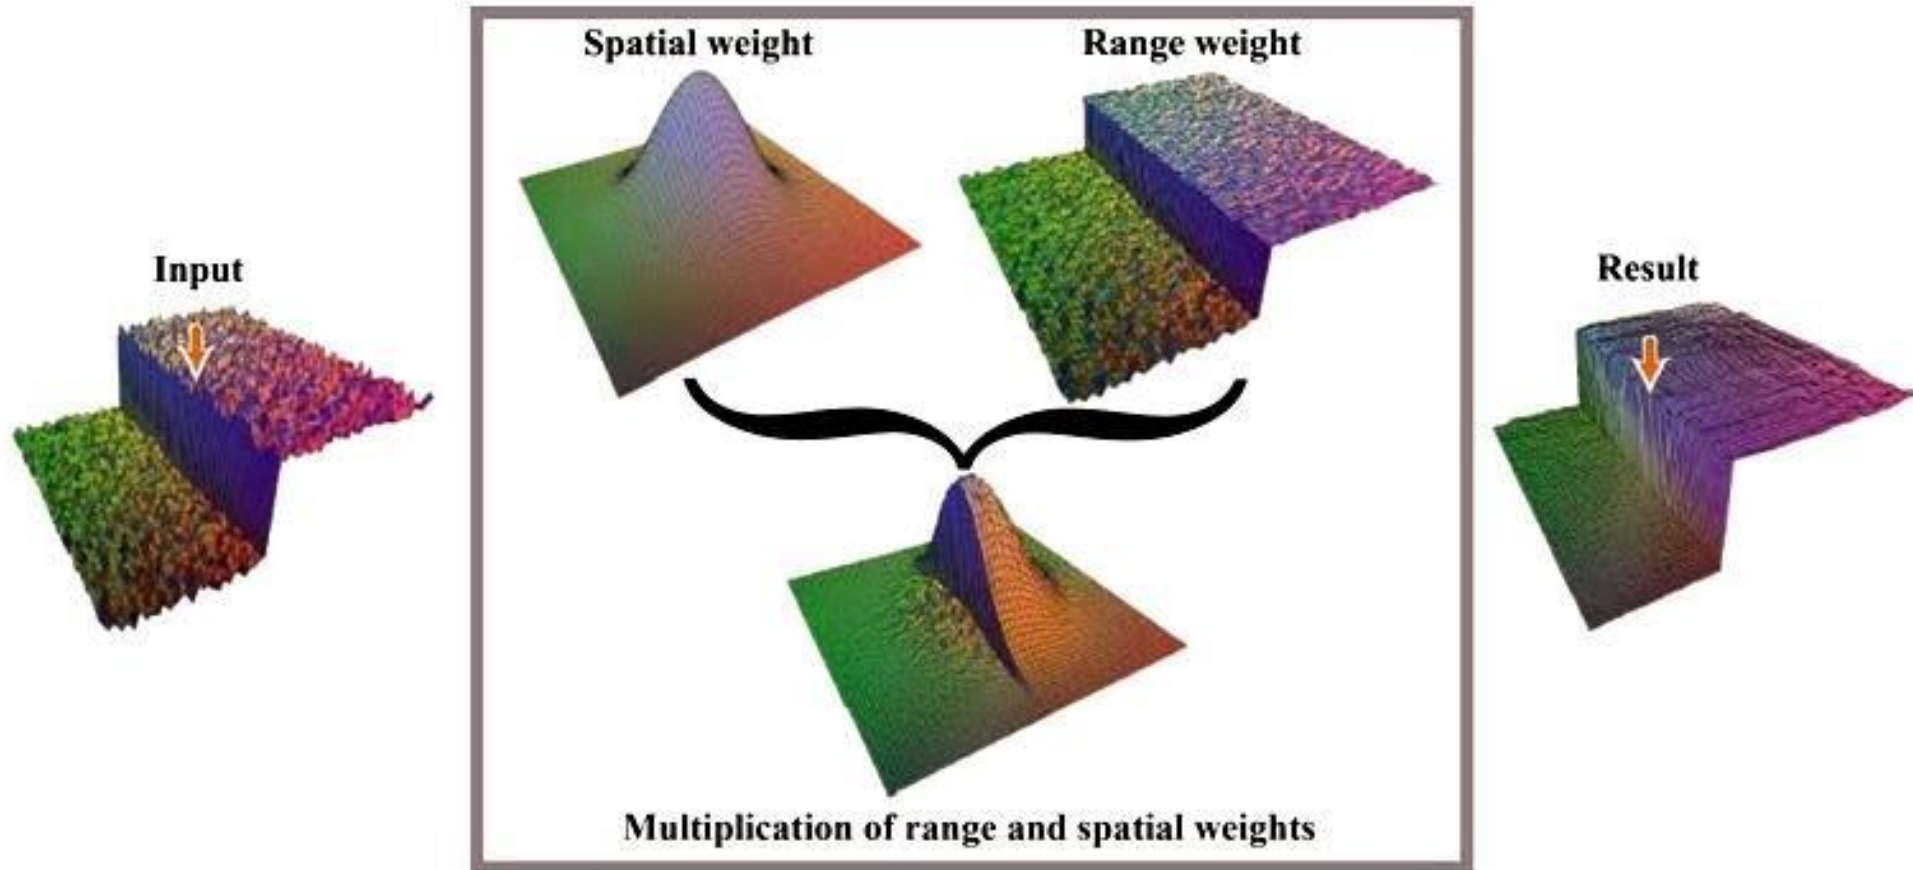
Median filter

1. Search for Neighbourhood values
2. Sort
3. Keep the median

25	151	84	34	62	132
1	224	71	188	178	71
71	153	120	111	238	184
65	61	14	15	26	226
58	144	72	41	94	191
152	66	153	184	18	225




Bilateral filter



Bilateral filter weights at the central pixel



Problem 2: Small holes / Salt noise



Morphology

- Morphology: A set of operations that expand or shrink object's shape
- 2 basic Morphological operations
 - Dilation: Expand the original shapes
 - Erosion: Shrink the original shape



After
Erosion



Original
object





After
Dilation



Morphology

- Process of dilation and erosion



Morphology

- Process of dilation and erosion


0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0
(a)				
0	1	1	1	0
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
0	1	1	1	0
(b)				
0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0
(c)				

Processes of dilation and erosion operations. (a) A 3×3 cross shape structural element applied to a 5×5 binary image. Each pixel value in the element represented the intensity. (b) Output image obtained from dilation. (c) Output image obtained from erosion.



Morphology

- **Kernel:** matrix/patterns that are used to process images




- **OpenCV function:**

`getStructuringElement(shape, size, anchor)`

```
// Morphology
Mat element = getStructuringElement( MORPH_RECT, Size(5,5), Point(-1,-1));
```



Morphology



- Element shape

Enumerator	
MORPH_RECT Python: cv.MORPH_RECT	a rectangular structuring element: $E_{ij} = 1$
MORPH_CROSS Python: cv.MORPH_CROSS	a cross-shaped structuring element: $E_{ij} = \begin{cases} 1 & \text{if } i = \text{anchor.y or } j = \text{anchor.x} \\ 0 & \text{otherwise} \end{cases}$
MORPH_ELLIPSE Python: cv.MORPH_ELLIPSE	an elliptic structuring element, that is, a filled ellipse inscribed into the rectangle Rect(0, 0, esize.width, 0.esize.height)



Morphology

- Single Operation

```
dilate(src,dst,element,anchor,iteration)
```

```
erode(src,dst,element,anchor,iteration)
```

- Advanced Transformation

```
morphologyEx(src,dst,operation,element,anchor,iteration) ;
```




Morphology

Enumerator	
MORPH_ERODE Python: cv.MORPH_ERODE	see erode
MORPH_DILATE Python: cv.MORPH_DILATE	see dilate
MORPH_OPEN Python: cv.MORPH_OPEN	an opening operation $\text{dst} = \text{open}(\text{src}, \text{element}) = \text{dilate}(\text{erode}(\text{src}, \text{element}))$
MORPH_CLOSE Python: cv.MORPH_CLOSE	a closing operation $\text{dst} = \text{close}(\text{src}, \text{element}) = \text{erode}(\text{dilate}(\text{src}, \text{element}))$
MORPH_GRADIENT Python: cv.MORPH_GRADIENT	a morphological gradient $\text{dst} = \text{morph_grad}(\text{src}, \text{element}) = \text{dilate}(\text{src}, \text{element}) - \text{erode}(\text{src}, \text{element})$
MORPH_TOPHAT Python: cv.MORPH_TOPHAT	"top hat" $\text{dst} = \text{tophat}(\text{src}, \text{element}) = \text{src} - \text{open}(\text{src}, \text{element})$
MORPH_BLACKHAT Python: cv.MORPH_BLACKHAT	"black hat" $\text{dst} = \text{blackhat}(\text{src}, \text{element}) = \text{close}(\text{src}, \text{element}) - \text{src}$
MORPH_HITMISS Python: cv.MORPH_HITMISS	"hit or miss" .- Only supported for CV_8UC1 binary images. A tutorial can be found in the documentation




Morphology

- Opening is defined as an **erosion followed by a dilation** using the *same structuring element* for both operations
- Closing is defined as a **dilation followed by an erosion** using the *same structuring element* for both operations.







Problem 3: Object cluster/occlusion

Cell stick to each other



Distance transform

- Distance transform convert matrix of pixel intensity values into distance map
 - Each pixel now show the shortest distance from itself to the nearest black area
 - Only apply for binary image



Distance transform

- Opencv Function:

`distanceTransform(src,dst,distanceType,markSize)`

Distance Type

type of the distance transformation to be applied

Mark Size

Size of the distance transform mask/kernel

```
// distance transform
Mat dist;
distanceTransform(gray,dst,DIST_L2,3);
// normalize to rang of [0 1]
normalize(dist, dist, 0, 1.0, NORM_MINMAX);
imshow("Distance Transform Image", dist);
```



Distance transform

- Distance type

Enumerator	
DIST_USER Python: cv.DIST_USER	User defined distance.
DIST_L1 Python: cv.DIST_L1	$\text{distance} = x1-x2 + y1-y2 $
DIST_L2 Python: cv.DIST_L2	the simple euclidean distance
DIST_C Python: cv.DIST_C	$\text{distance} = \max(x1-x2 , y1-y2)$
DIST_L12 Python: cv.DIST_L12	L1-L2 metric: $\text{distance} = 2(\sqrt{1+x*x/2}) - 1$
DIST_FAIR Python: cv.DIST_FAIR	$\text{distance} = c^2(x /c - \log(1+ x /c))$, $c = 1.3998$
DIST_WELSCH Python: cv.DIST_WELSCH	$\text{distance} = c^2/2(1 - \exp(-(x/c)^2))$, $c = 2.9846$
DIST_HUBER Python: cv.DIST_HUBER	$\text{distance} = x < c ? x^2/2 : c(x -c/2)$, $c=1.345$



Distance transform

- Mask size

+b	+a	+b
+a	0	+a
+b	+a	+b


3 x 3 mask

-	+c	-	+c	-
+c	+b	+a	+b	+c
-	+a	0	+a	-
+c	+b	+a	+b	+c
-	+c	-	+c	-


5 x 5 mask





Distance transform





Distance Transform



- Normalized to [0 1]
- Threshold



Watershed



https://docs.opencv.org/3.4/d2/dbd/tutorial_distance_transform.html




Watershed

- Opencv Function:

`watershed(image, markers)`



`image`



Markers
(before)



Result

